

# Analyze\_ab\_test\_results\_notebook

November 23, 2021

## 1 Analyze A/B Test Results

This project will assure you have mastered the subjects covered in the statistics lessons. We have organized the current notebook into the following sections:

- Section ??
- Section ??
- Section ??
- Section ??
- Section ??

Specific programming tasks are marked with a **ToDo** tag.

## Introduction

A/B tests are very commonly performed by data analysts and data scientists. For this project, you will be working to understand the results of an A/B test run by an e-commerce website. Your goal is to work through this notebook to help the company understand if they should: - Implement the new webpage, - Keep the old webpage, or - Perhaps run the experiment longer to make their decision.

Each **ToDo** task below has an associated quiz present in the classroom. Though the classroom quizzes are **not necessary** to complete the project, they help ensure you are on the right track as you work through the project, and you can feel more confident in your final submission meeting the [rubric](#) specification.

## Part I - Probability

To get started, let's import our libraries.

```
In [1]: import pandas as pd
import numpy as np
import datetime as dt
import random
import matplotlib.pyplot as plt
%matplotlib inline
#We are setting the seed to assure you get the same answers on quizzes as we set up
random.seed(42)
```

### 1.0.1 ToDo 1.1

Now, read in the `ab_data.csv` data. Store it in `df`.

- a. Read in the dataset from the `ab_data.csv` file and take a look at the top few rows here:

```
In [2]: # importing 'ab_data.csv' data file
df = pd.read_csv ("ab_data.csv")
# Display the top five rows from the loaded data file
df.head()
```

```
Out[2]:
```

	user_id	timestamp	group	landing_page	converted
0	851104	2017-01-21 22:11:48.556739	control	old_page	0
1	804228	2017-01-12 08:01:45.159739	control	old_page	0
2	661590	2017-01-11 16:55:06.154213	treatment	new_page	0
3	853541	2017-01-08 18:28:03.143765	treatment	new_page	0
4	864975	2017-01-21 01:52:26.210827	control	old_page	1

b. Use the cell below to find the number of rows in the dataset.

```
In [3]: # Display the number of rows
df.shape[0]
```

```
Out[3]: 294478
```

c. The number of unique users in the dataset.

```
In [4]: # Display the number of unique users
df['user_id'].nunique()
```

```
Out[4]: 290584
```

d. The proportion of users converted.

```
In [5]: # The proportion of users converted
df['converted'].mean()
# Alternative code
df['converted'].sum()/df.shape[0]
```

```
Out[5]: 0.11965919355605512
```

e. The number of times when the "group" is treatment but "landing\_page" is not a new\_page.

```
In [6]: # Display the number of times the "group" is treatment but "landing_page" is not a new_page
df.query('group == "treatment" & landing_page != "new_page").shape[0]
```

```
Out[6]: 1965
```

f. Do any of the rows have missing values?

```
In [7]: # Display the number of rows having null values
df.shape[0] - df.dropna().shape[0]
```

```
Out[7]: 0
```

No rows with missing data in the dataset

## 1.0.2 ToDo 1.2

The rows where treatment does not match with new\_page or control does not match with old\_page, we cannot be sure if such rows truly received the new or old webpage.

a. Now use the answer to the quiz to create a new dataset that meets the specifications from the quiz. Store your new dataframe in **df2**.

```
In [8]: # Remove the inaccurate rows where the 'group' and 'landing_page' columns don't match, a
df2 = df.drop(df.query('group == "treatment" & landing_page != "new_page" | group == "co
df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 290585 entries, 0 to 294477
Data columns (total 5 columns):
user_id          290585 non-null int64
timestamp        290585 non-null object
group            290585 non-null object
landing_page     290585 non-null object
converted        290585 non-null int64
dtypes: int64(2), object(3)
memory usage: 13.3+ MB
```

```
In [9]: # Double Check all of the incorrect rows were removed from df2 -
# Output of the statement below should be 0
df2[((df2['group'] == 'treatment') == (df2['landing_page'] == 'new_page')) == False].sha
```

```
Out[9]: 0
```

## 1.0.3 ToDo 1.3

Use **df2** and the cells below to answer questions for **Part 1.3**.

a. How many unique **user\_ids** are in **df2**?

```
In [10]: # The number of unique users in df2 from 'user_id' column
df2['user_id'].nunique()
```

```
Out[10]: 290584
```

b. There is one **user\_id** repeated in **df2**. What is it?

```
In [11]: # Display the duplicated users from 'user_id' column
df2[df2['user_id'].duplicated()]['user_id']
```

```
Out[11]: 2893      773192
Name: user_id, dtype: int64
```

c. Display the rows for the duplicate **user\_id**?

```
In [12]: # Display a row with duplicated 'user_id' information
df2[df2['user_id'].duplicated()]
```

```
Out[12]:
```

	user_id	timestamp	group	landing_page	converted	
	2893	773192	2017-01-14 02:55:59.590927	treatment	new_page	0

d. Remove **one** of the rows with a duplicate **user\_id**, from the **df2** dataframe.

```
In [13]: # Remove one of the rows with a duplicate user_id..
# Hint: The dataframe.drop_duplicates() may not work in this case because the rows with
df2 = df2.drop(df2[df2.duplicated(['user_id'])].index)
# Check again if the row with a duplicate user_id is deleted or not
df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 290584 entries, 0 to 294477
Data columns (total 5 columns):
user_id      290584 non-null int64
timestamp    290584 non-null object
group        290584 non-null object
landing_page 290584 non-null object
converted    290584 non-null int64
dtypes: int64(2), object(3)
memory usage: 13.3+ MB
```

## 1.0.4 ToDo 1.4

Use **df2** and the cells below to answer questions for **Part 1.4**.

a. What is the probability of an individual converting regardless of the page they receive?

```
In [14]: # Display the probability of an individual converting
#This is also the overall "converted" success rate in the population (p-population).
conversion_rate = df2['converted'].mean()
conversion_rate
```

```
Out[14]: 0.11959708724499628
```

b. Given that an individual was in the control group, what is the probability they converted?

```
In [15]: # Display the probability of an individual converting from the "control" group
control_propability = df2.query('group == "control" & converted == 1').shape[0] / df2.q
control_propability
```

```
Out[15]: 0.1203863045004612
```

c. Given that an individual was in the treatment group, what is the probability they converted?

```
In [16]: # Display the probability of an individual converting from the "treatment" group
treatment_propability = df2.query('group == "treatment" & converted == 1').shape[0] / d
treatment_propability
```

```
Out[16]: 0.11880806551510564
```

```
In [17]: # Calculate the actual difference (obs_diff) between the conversion rates for the two g
        obs_diff = treatment_propability - control_propability
        obs_diff
```

```
Out[17]: -0.0015782389853555567
```

d. What is the probability that an individual received the new page?

```
In [18]: # Display the probability of an individual received the new page
        new_page_propability = df2.query('landing_page == "new_page").count()[0] / df.shape[0]
        new_page_propability
```

```
Out[18]: 0.49344942576355449
```

e. Consider your results from parts (a) through (d) above, and explain below whether the new treatment group users lead to more conversions.

The difference in **conversion rate** between the control and the treatment groups is **15%**; Such a minor difference in the conversion rate does not indicate any sufficient evidence to suggest: The new treatment page leads to more conversions.

## Part II - A/B Test

Since a timestamp is associated with each event, you could run a hypothesis test continuously as long as you observe the events.

However, then the hard questions would be: - Do you stop as soon as one page is considered significantly better than another or does it need to happen consistently for a certain amount of time?

- How long do you run to render a decision that neither page is better than another?

These questions are the difficult parts associated with A/B tests in general.

### 1.0.5 ToDo 2.1

For now, consider you need to make the decision just based on all the data provided.

If you want to assume that the old page is better unless the new page proves to be definitely better at a Type I error rate of 5%, what should be your null and alternative hypotheses ( $H_0$  and  $H_1$ )?

You can state your hypothesis in terms of words or in terms of  $p_{old}$  and  $p_{new}$ , which are the "converted" probability (or rate) for the old and new pages respectively.

$H_0 : p_{new} \leq p_{old}$   $H_1 : p_{new} > p_{old}$

### 1.0.6 ToDo 2.2 - Null Hypothesis $H_0$ Testing

Under the null hypothesis  $H_0$ , assume that  $p_{new}$  and  $p_{old}$  are equal. Furthermore, assume that  $p_{new}$  and  $p_{old}$  both are equal to the **converted** success rate in the df2 data regardless of the page. So, our assumption is:

$p_{new} = p_{old} = p_{population}$

Use the cells below to provide the necessary parts of this simulation.

a. What is the **conversion rate** for  $p_{new}$  under the null hypothesis?

```
In [19]: conversion_rate
```

```
Out[19]: 0.11959708724499628
```

b. What is the **conversion rate** for  $p_{old}$  under the null hypothesis?

```
In [20]: conversion_rate
```

```
Out[20]: 0.11959708724499628
```

c. What is  $n_{new}$ , the number of individuals in the treatment group? *Hint:* The treatment group users are shown the new page.

```
In [21]: n_new = df2.query('group == "treatment"').shape[0]
         n_new
```

```
Out[21]: 145310
```

d. What is  $n_{old}$ , the number of individuals in the control group?

```
In [22]: n_old = df2.query('group == "control"').shape[0]
         n_old
```

```
Out[22]: 145274
```

e. **Simulate Sample for the treatment Group** Simulate  $n_{new}$  transactions with a conversion rate of  $p_{new}$  under the null hypothesis. Store these  $n_{new}$  1's and 0's in the `new_page_converted` numpy array.

```
In [23]: # Simulate a Sample for the treatment Group
         # new_page_converted = np.random.choice(2, n_new, replace = True, p=[(1-conversion_rate)
         new_page_converted = np.random.choice(2, n_new, conversion_rate)
         new_page_converted
```

```
Out[23]: array([0, 1, 1, ..., 1, 0, 0])
```

f. **Simulate Sample for the control Group** Simulate  $n_{old}$  transactions with a conversion rate of  $p_{old}$  under the null hypothesis. Store these  $n_{old}$  1's and 0's in the `old_page_converted` numpy array.

```
In [24]: # Simulate a Sample for the control Group
         old_page_converted = np.random.choice(2, n_old, conversion_rate)
         old_page_converted
```

```
Out[24]: array([1, 0, 0, ..., 0, 1, 0])
```

g. Find the difference in the "converted" probability ( $p'_{new} - p'_{old}$ ) for your simulated samples from the parts (e) and (f) above.

```
In [25]: p_new = new_page_converted.mean()
         p_old = old_page_converted.mean()
         converted_probability_diff = p_new - p_old
         converted_probability_diff
```

```
Out[25]: -0.00293901927266077
```


**h. Sampling distribution** Re-create `new_page_converted` and `old_page_converted` and find the  $(p'_{new} - p'_{old})$  value 10,000 times using the same simulation process you used in parts (a) through (g) above. Store all  $(p'_{new} - p'_{old})$  values in a NumPy array called `p_diffs`.

```
In [26]: # Sampling distribution
p_diffs = []
p_new = conversion_rate
p_old = conversion_rate
new_simulation = np.random.binomial(n_new, p_new, 10000)/n_new
old_simulation = np.random.binomial(n_old, p_old, 10000)/n_old
p_diffs = new_simulation - old_simulation
p_diffs
```

```
Out[26]: array([-1.11698013e-03, -1.59885551e-03,  3.69532869e-04, ...,
               -9.82889201e-05, -1.57145776e-03, -7.45695087e-04])
```


**i. Histogram** Plot a histogram of the `p_diffs`. Does this plot look like what you expected? Use the matching problem in the classroom to assure you fully understand what was computed here. Also, use `plt.axvline()` method to mark the actual difference observed in the `df2` data (recall `obs_diff`), in the chart.

```
In [27]: plt.hist(p_diffs)
plt.show()
```



```
In [28]: h_null = np.random.normal(0, np.std(p_diffs),10000)

plt.hist(h_null);
plt.axvline(converted_probability_diff, color = 'red');
plt.axvline(-converted_probability_diff, color = 'red');
```



j. What proportion of the **p\_diffs** are greater than the actual difference observed in the df2 data?

```
In [29]: (p_diffs > obs_diff).mean()
```

```
Out[29]: 0.90620000000000001
```

k. Please explain in words what you have just computed in part j above.

- What is this value called in scientific studies?
- What does this value signify in terms of whether or not there is a difference between the new and old pages? *Hint:* Compare the value above with the "Type I error rate (0.05)".

- What we computed in point (j) was the p-value, which is the probability that we will observe this statistic, in the case that the null hypothesis is true.
- Since the p-value is higher than the **Type I error rate** of (5%), **we fail to reject the null hypothesis** and the treatment page does not have higher conversion rates than the control page on a statistical basis.

l. **Using Built-in Methods for Hypothesis Testing** We could also use a built-in to achieve similar results. Though using the built-in might be easier to code, the above portions are a walk-through of the ideas that are critical to correctly thinking about statistical significance.



```
In [30]: import statsmodels.api as sm
# number of conversions with the old_page
convert_old = df.query('group == "control" & converted == 1').shape[0]
# number of conversions with the new_page
convert_new = df2.query('group == "treatment" & converted == 1').shape[0]
# number of individuals who were shown the old_page
n_old = df2.query("landing_page == 'old_page'").shape[0]
# number of individuals who received new_page
n_new = df2.query("landing_page == 'new_page'").shape[0]
# Testing our results
convert_old, convert_new, n_old, n_new
```

```
/opt/conda/lib/python3.6/site-packages/statsmodels/compat/pandas.py:56: FutureWarning: The pandas
from pandas.core import datetools
```

```
Out[30]: (17723, 17264, 145274, 145310)
```

m. Now use `sm.stats.proportions_ztest()` to compute your test statistic and p-value. [Here](#) is a helpful link on using the built in.

```
In [31]: import statsmodels.api as sm
# ToDo: Complete the sm.stats.proportions_ztest() method arguments
z_score, p_value = sm.stats.proportions_ztest([convert_new, convert_old], [n_new, n_old])
z_score, p_value
```

```
Out[31]: (-2.6411875940931551, 0.99586920197633411)
```

n. What do the z-score and p-value you computed in the previous question mean for the conversion rates of the old and new pages? Do they agree with the findings in parts j. and k.?

The z-score and p-value we computed from the test statistic suggests that the conversion rates for both old and new pages are not statistically different from one another, the p-value (0.09) still larger than the **type I error rate** (0.05) which also affirms **not rejecting the null hypothesis**.

### Part III - A regression approach

### 1.0.7 ToDo 3.1

In this final part, you will see that the result you achieved in the A/B test in Part II above can also be achieved by performing regression.

a. Since each row in the `df2` data is either a conversion or no conversion, what type of regression should you be performing in this case?

A logistic regression since we are looking at *binary variables*.

b. The goal is to use **statsmodels** library to fit the regression model you specified in part a. above to see if there is a significant difference in conversion based on the page-type a customer receives. However, you first need to create the following two columns in the `df2` dataframe: 1. `intercept` - It should be 1 in the entire column. 2. `ab_page` - It's a dummy variable column, having a value 1 when an individual receives the **treatment**, otherwise 0.

```
In [32]: df2['intercept'] = 1
df2['ab_page'] = pd.get_dummies(df2['group'])['treatment']
df2.head()
```

```

Out[32]:   user_id      timestamp      group landing_page  converted  \
0   851104  2017-01-21 22:11:48.556739   control   old_page         0
1   804228  2017-01-12 08:01:45.159739   control   old_page         0
2   661590  2017-01-11 16:55:06.154213  treatment   new_page         0
3   853541  2017-01-08 18:28:03.143765  treatment   new_page         0
4   864975  2017-01-21 01:52:26.210827   control   old_page         1

      intercept  ab_page
0             1         0
1             1         0
2             1         1
3             1         1
4             1         0

```

c. Use **statsmodels** to instantiate your regression model on the two columns you created in part (b). above, then fit the model to predict whether or not an individual converts.

```

In [33]: my_reg_model = sm.Logit(df2['converted'], df2[['intercept', 'ab_page']])
        result = my_reg_model.fit()

```

```

Optimization terminated successfully.
Current function value: 0.366118
Iterations 6

```

d. Provide the summary of your model below, and use it as necessary to answer the following questions.

```

In [34]: result.summary2()

```

```

Out[34]: <class 'statsmodels.iolib.summary2.Summary'>
        """
                Results: Logit
        =====
Model:                Logit                No. Iterations:    6.0000
Dependent Variable:  converted                Pseudo R-squared:  0.000
Date:                2021-11-23 18:58    AIC:                212780.3502
No. Observations:    290584                BIC:                212801.5095
Df Model:            1                    Log-Likelihood:    -1.0639e+05
Df Residuals:        290582                LL-Null:           -1.0639e+05
Converged:           1.0000                Scale:            1.0000
-----
                Coef.   Std.Err.    z      P>|z|    [0.025   0.975]
-----
intercept    -1.9888    0.0081  -246.6690  0.0000   -2.0046   -1.9730
ab_page      -0.0150    0.0114   -1.3109  0.1899   -0.0374    0.0074
=====
        """

```

```
In [35]: np.exp(result.params)
```

```
Out[35]: intercept    0.136863  
         ab_page      0.985123  
         dtype: float64
```

```
In [36]: 1/np.exp(result.params)
```

```
Out[36]: intercept    7.306593  
         ab_page      1.015102  
         dtype: float64
```

e. What is the p-value associated with **ab\_page**? Why does it differ from the value you found in **Part II**?

- Both cases p-value are higher than the accepted **Type I error rate (0.05)**, the results do not support the alternative hypothesis sufficiently.
- The p-value is different in **part II (0.91)** from **part III (0.19)** because of using different null and alternative hypothesis here.  $H_0 : p_{new} - p_{old} = 0$   $H_1 : p_{new} - p_{old} \neq 0$

f. Now, you are considering other things that might influence whether or not an individual converts. Discuss why it is a good idea to consider other factors to add into your regression model. Are there any disadvantages to adding additional terms into your regression model?

- It is a good idea to incorporate more factors that could potentially influence the conversion rate into the model, like users (**old & new**), this data includes a single factor influencing the conversion rate page type.
- The disadvantages comes from the data itself, using incomplete data leads to erroneous conclusions.

g. **Adding countries** Now along with testing if the conversion rate changes for different pages, also add an effect based on which country a user lives in.

1. You will need to read in the **countries.csv** dataset and merge together your df2 datasets on the appropriate rows. You call the resulting dataframe df\_merged. [Here](#) are the docs for joining tables.
2. Does it appear that country had an impact on conversion? To answer this question, consider the three unique values, ['UK', 'US', 'CA'], in the country column. Create dummy variables for these country columns.

Provide the statistical output as well as a written response to answer this question.

```
In [37]: # Reading the countries.csv and viewing the first 5 rows of data  
         countries = pd.read_csv("countries.csv")  
         countries.head()
```

```
Out[37]:   user_id country  
0    834778      UK  
1    928468      US  
2    822059      UK  
3    711597      UK  
4    710616      UK
```

```
In [38]: # Joining countries and df2 dataframes into df_merged and viewing the first 5 rows of a
df_merged = countries.merge(df2)
df_merged.head()
```

```
Out[38]:
```

	user_id	country	timestamp	group	landing_page	\
0	834778	UK	2017-01-14 23:08:43.304998	control	old_page	
1	928468	US	2017-01-23 14:44:16.387854	treatment	new_page	
2	822059	UK	2017-01-16 14:04:14.719771	treatment	new_page	
3	711597	UK	2017-01-22 03:14:24.763511	control	old_page	
4	710616	UK	2017-01-16 13:14:44.000513	treatment	new_page	

	converted	intercept	ab_page
0	0	1	0
1	0	1	1
2	1	1	1
3	0	1	0
4	0	1	1

```
In [39]: # Create the necessary dummy variables
df_merged[['UK', 'US', 'CA']] = pd.get_dummies(df_merged['country'])
df_merged.head()
```

```
Out[39]:
```

	user_id	country	timestamp	group	landing_page	\
0	834778	UK	2017-01-14 23:08:43.304998	control	old_page	
1	928468	US	2017-01-23 14:44:16.387854	treatment	new_page	
2	822059	UK	2017-01-16 14:04:14.719771	treatment	new_page	
3	711597	UK	2017-01-22 03:14:24.763511	control	old_page	
4	710616	UK	2017-01-16 13:14:44.000513	treatment	new_page	

	converted	intercept	ab_page	UK	US	CA
0	0	1	0	0	1	0
1	0	1	1	0	0	1
2	1	1	1	0	1	0
3	0	1	0	0	1	0
4	0	1	1	0	1	0

```
In [40]: # Fitting our model, and summarize the results
country_reg_model = sm.Logit(df_merged['converted'], df_merged[['intercept', 'ab_page',
results = country_reg_model.fit()
results.summary2()
```

```
Optimization terminated successfully.
Current function value: 0.366113
Iterations 6
```

```
Out[40]: <class 'statsmodels.iolib.summary2.Summary'>
"""
```

Results: Logit

```

=====
Model:                Logit                No. Iterations:    6.0000
Dependent Variable: converted                Pseudo R-squared: 0.000
Date:                2021-11-23 18:58 AIC:                212781.1253
No. Observations:    290584                BIC:                212823.4439
Df Model:            3                    Log-Likelihood:    -1.0639e+05
Df Residuals:        290580                LL-Null:            -1.0639e+05
Converged:            1.0000                Scale:            1.0000
-----
                Coef.    Std.Err.    z        P>|z|    [0.025    0.975]
-----
intercept      -2.0300     0.0266   -76.2488  0.0000   -2.0822   -1.9778
ab_page        -0.0149     0.0114   -1.3069  0.1912   -0.0374    0.0075
US              0.0506     0.0284    1.7835  0.0745   -0.0050    0.1063
CA              0.0408     0.0269    1.5161  0.1295   -0.0119    0.0934
=====
"""

```

```
In [41]: np.exp(results.params)
```

```

Out[41]: intercept    0.131332
         ab_page      0.985168
         US           1.051944
         CA           1.041599
         dtype: float64

```

```
In [42]: 1/np.exp(results.params)
```

```

Out[42]: intercept    7.614303
         ab_page      1.015056
         US           0.950621
         CA           0.960062
         dtype: float64

```

**Statistically:** - The country of residence had no impact on conversion rate **US 1.05** and **CA 1.04** are higher than **UK** (Baseline). - The difference between US and UK (baseline) =  $(1.051944 - 1) * 100 = 5.1944\%$  - The difference between CA and UK (baseline) =  $(1.041599 - 1) * 100 = 4.1599\%$

**Practically:** - Countries are irrelevant to include in this analysis.

**h. Fit your model and obtain the results** Though you have now looked at the individual factors of country and page on conversion, we would now like to look at an interaction between page and country to see if there are significant effects on conversion. **Create the necessary additional columns, and fit the new model.**

Provide the summary results (statistical output), and your conclusions (written response) based on the results.

```

In [43]: df_merged["ab_UK"] = df_merged["ab_page"] * df_merged["UK"]
         df_merged["ab_CA"] = df_merged["ab_page"] * df_merged["CA"]

```

```
In [44]: legit_model = sm.Logit(df_merged['converted'], df_merged[['intercept', 'ab_page', 'US'],
    legit_fit = legit_model.fit())
```

```
Optimization terminated successfully.
    Current function value: 0.366109
    Iterations 6
```

```
In [45]: legit_fit.summary2()
```

```
Out[45]: <class 'statsmodels.iolib.summary2.Summary'>
        """
```

```

                                Results: Logit
=====
Model:                Logit                No. Iterations:    6.0000
Dependent Variable: converted                Pseudo R-squared: 0.000
Date:                2021-11-23 18:58 AIC:                212782.6602
No. Observations:    290584                BIC:                212846.1381
Df Model:            5                    Log-Likelihood:    -1.0639e+05
Df Residuals:        290578                LL-Null:            -1.0639e+05
Converged:            1.0000                Scale:            1.0000
-----
                Coef.    Std.Err.    z      P>|z|    [0.025    0.975]
-----
intercept      -2.0040    0.0364   -55.0077  0.0000   -2.0754   -1.9326
ab_page         0.0108    0.0228    0.4749  0.6349   -0.0339    0.0555
US              0.0118    0.0398    0.2957  0.7674   -0.0663    0.0899
CA              0.0175    0.0377    0.4652  0.6418   -0.0563    0.0914
ab_UK          -0.0783    0.0568   -1.3783  0.1681   -0.1896    0.0330
ab_CA          -0.0314    0.0266   -1.1807  0.2377   -0.0835    0.0207
=====
        """
```

```
In [46]: np.exp(legit_fit.params)
```

```
Out[46]: intercept    0.134794
         ab_page      1.010893
         US           1.011854
         CA           1.017682
         ab_UK        0.924703
         ab_CA        0.969090
         dtype: float64
```

```
In [47]: 1/np.exp(legit_fit.params)
```

```
Out[47]: intercept    7.418713
         ab_page      0.989224
         US           0.988285
```

```
CA          0.982625
ab_UK       1.081428
ab_CA       1.031896
dtype: float64
```

```
In [48]: # Checning duration of the experiment
df_merged['timestamp'] = df_merged['timestamp'].apply(lambda x : pd.to_datetime(str(x)))
df_merged['dates'] = df_merged['timestamp'].dt.date
```

```
In [49]: ((df_merged['dates'].max() - df_merged['dates'].min()), 'days')
```

```
Out[49]: (datetime.timedelta(22), 'days')
```

**Statistically:** - The conversion by country has no statistical significance since all p\_values are large in all cases we failed to reject null.

**Practically:** - The impact of conversion by country is with **no significance**, the conversion rate for both US and CA is higher than UK.

### Final Conclusion

- Countries conversion rate was irrelevant to the case.
- Duration of the experiment was 22 days.
- There is 15% difference in the conversion rate between the control and treatment groups from the dataset itself, which imposed almost no effect to influence the decision to switch to the new page.
- We failed to reject the Null in all the 3 cases, **Descriptive analysis, A/B Testing, and Regression Modelling.**

Run the code cell below to create a .html version of this notebook in the workspace. If it worked correctly, you should get a return code of 0

```
In [50]: from subprocess import call
         call(['python', '-m', 'nbconvert', 'Analyze_ab_test_results_notebook.ipynb'])
```

```
Out[50]: 0
```